

# Contents

<b>Scenario: Replicate Data Across Regions</b>	<b>1</b>
Purpose . . . . .	1
Summary . . . . .	1
Replication Strategies . . . . .	1
Consistency Models . . . . .	2
CAP Theorem . . . . .	2
PACELC Theorem . . . . .	2
Conflict Resolution . . . . .	2
Failover & Disaster Recovery . . . . .	3
Network & Latency Considerations . . . . .	3
Security & Compliance . . . . .	3
Metrics for Success . . . . .	3

## Scenario: Replicate Data Across Regions

*Describe your approach to cross-region data replication, consistency, and failover.*

### Purpose

- Tests your understanding of distributed systems and data consistency
  - Evaluates your ability to balance latency, availability, and durability
  - Assesses your approach to conflict resolution and disaster recovery
  - Checks your awareness of trade-offs in global architectures
- 

### Summary

Choose async replication for low-latency writes. Use quorum-based reads for stronger consistency. Resolve conflicts using timestamps or version vectors.

---

### Replication Strategies

- **Asynchronous Replication:**
    - Primary region handles writes, replicates changes to secondary regions in the background
    - Pros: Low write latency, high availability
    - Cons: Risk of data loss on failover, eventual consistency
  - **Synchronous Replication:**
    - Writes are committed only after all (or a quorum of) regions acknowledge
    - Pros: Stronger consistency, no data loss on failover
    - Cons: Higher write latency, lower availability during network partitions
  - **Hybrid/Quorum-Based:**
    - Use a quorum of regions for reads/writes (e.g., majority must agree)
    - Balances consistency and availability
-

## Consistency Models

- **Eventual Consistency:** All regions will converge to the same state, but reads may be stale
  - **Strong Consistency:** Reads always return the latest committed write (requires synchronous/quorum replication)
  - **Tunable Consistency:** Allow clients to choose consistency level per operation (e.g., Cassandra, DynamoDB)
- 

## CAP Theorem

The CAP theorem states that in a distributed data system, you can only guarantee two out of the following three properties in the presence of a network partition: - **Consistency (C):** Every read receives the most recent write or an error. - **Availability (A):** Every request receives a (non-error) response, regardless of the state of other nodes. - **Partition Tolerance (P):** The system continues to operate despite network partitions (communication breakdowns between nodes/regions).

**Key clarification:** - When there is no network partition, a system can provide both Consistency and Availability. - When a partition occurs (which is inevitable in distributed, cross-region systems), the system must choose to either: - Remain **consistent** (some requests may be denied to prevent stale reads), or - Remain **available** (all requests are served, but some may return stale data).

**In cross-region replication:** - Partition tolerance is required due to the nature of wide-area networks. - The trade-off between Consistency and Availability only arises during a partition event.

---

## PACELC Theorem

PACELC extends CAP by considering system behavior not only during partitions, but also under normal operation: - **If there is a Partition (P):** choose between Availability (A) and Consistency (C) — just like CAP. - **Else (E), when the system is running normally:** choose between Latency (L) and Consistency (C).

**In practice:** - Systems like DynamoDB, Cassandra, and Cosmos DB allow tuning this trade-off. - You may choose lower latency (faster responses) at the cost of weaker consistency, or stronger consistency with higher latency, even when there is no partition.

**Summary Table:** | Theorem | Partition? | Trade-off | | CAP | Yes | C vs. A | | PACELC | Yes | C vs. A | | PACELC | No | C vs. L |

---

## Conflict Resolution

- **Last Write Wins (timestamp-based):** Simple, but may lose updates
  - **Version Vectors:** Track causality and resolve conflicts more accurately
  - **Custom Merge Logic:** For complex data types (e.g., CRDTs)
-

## Failover & Disaster Recovery

- **Automated Failover:** Detect region failure and promote a new primary
  - **Data Reconciliation:** Sync missed updates after recovery
  - **Backup & Restore:** Regular cross-region backups for disaster recovery
- 

## Network & Latency Considerations

- Use geo-DNS or global load balancers to route users to the nearest healthy region
  - Minimize cross-region traffic for latency-sensitive operations
  - Compress and batch replication traffic
- 

## Security & Compliance

- Encrypt data in transit and at rest across all regions
  - Ensure compliance with data residency and privacy laws (e.g., GDPR)
- 

## Metrics for Success

- Replication lag (seconds behind primary)
- Data consistency error rate
- Recovery time objective (RTO) and recovery point objective (RPO)
- User-perceived latency and availability